



# Attention in Deep Learning



# Visual Attention & Attention in Text

Attention is **the ability to choose and concentrate on relevant stimuli**. In neural networks, attention is a technique that mimics cognitive attention. The effect enhances some parts of the input data while diminishing other parts. In other words, **attention is the way how we pay visual attention to different regions** of an image or correlate words in one sentence.



high attention

low attention

She is eating a green apple.

# Attention in Deep Learning

**Machine Translation (MT)** is the task of translating a sentence  $x$  from one language (the **source language**) to a sentence  $y$  in another language (the **target language**).

$x$ : *L'homme est né libre, et partout il est dans les fers*




$y$ : *Man is born free, but everywhere he is in chains*



Attention was initially designed in the context of Neural Machine Translation using Seq2Seq Models


**Alignment, i.e. word-level correspondence between source sentence  $x$  and target sentence  $y$**



# What is Alignment

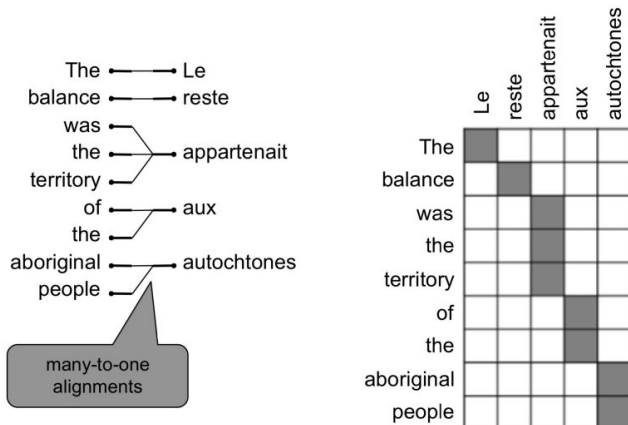
Alignment is the **correspondence between particular words** in the translated sentence pair.

- **Typological differences** between languages lead to complicated alignments!
- Note: Some words have **no counterpart**

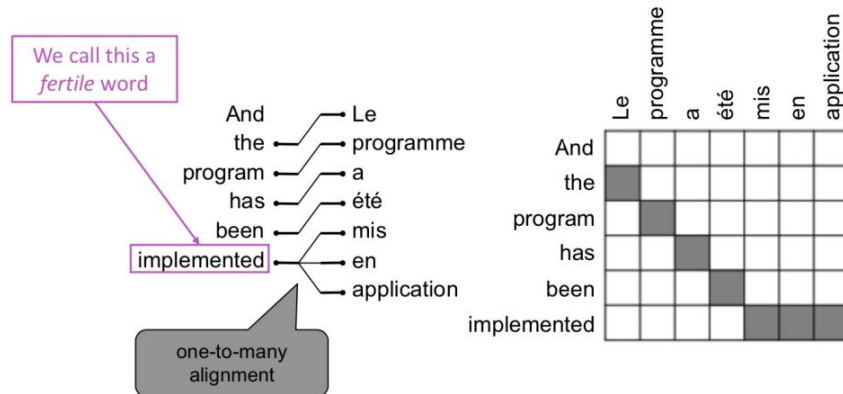


# Alignment is Complex

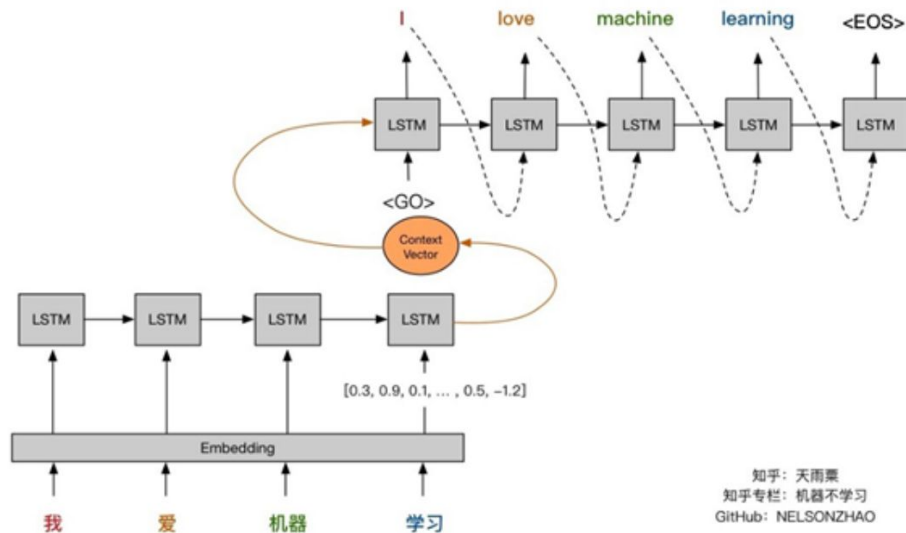
Alignment can be **many-to-one**



Alignment can be **one-to-many**



# Encoder-Decoder in seq2seq



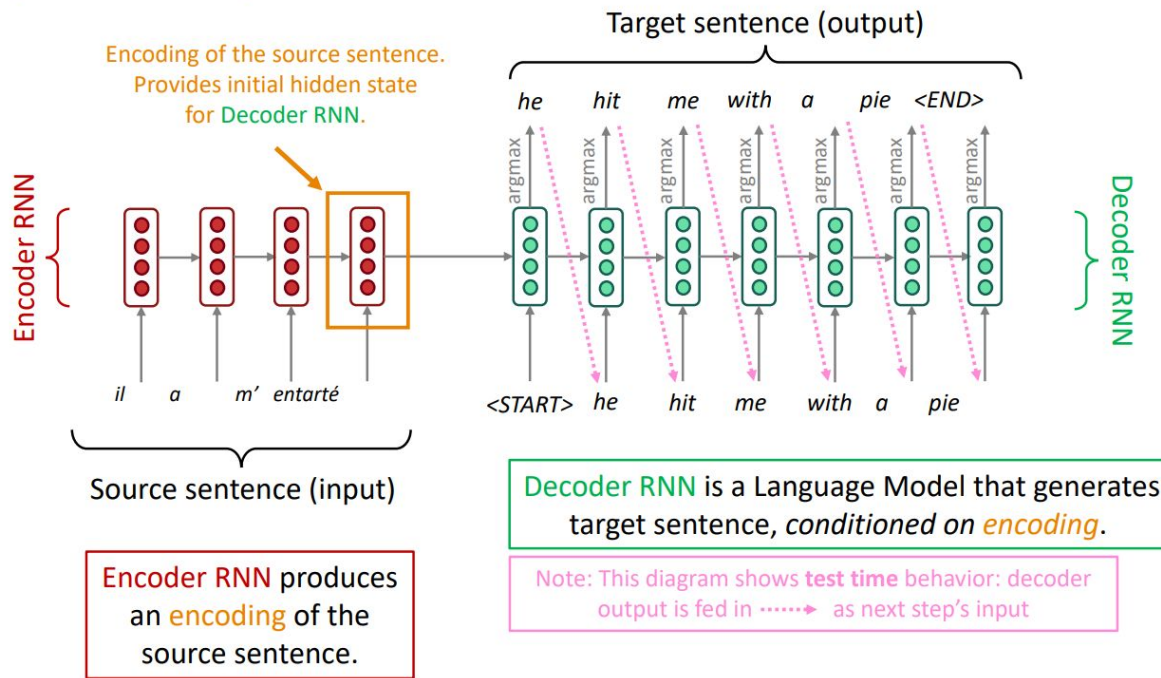
知乎: 天雨粟  
知乎专栏: 机器不学习  
GitHub: NELSONZHAO

知乎 @天雨粟

Encoder-Decoder with simple fixed context vector

# Neural Machine Translation aka seq2seq

## The sequence-to-sequence model

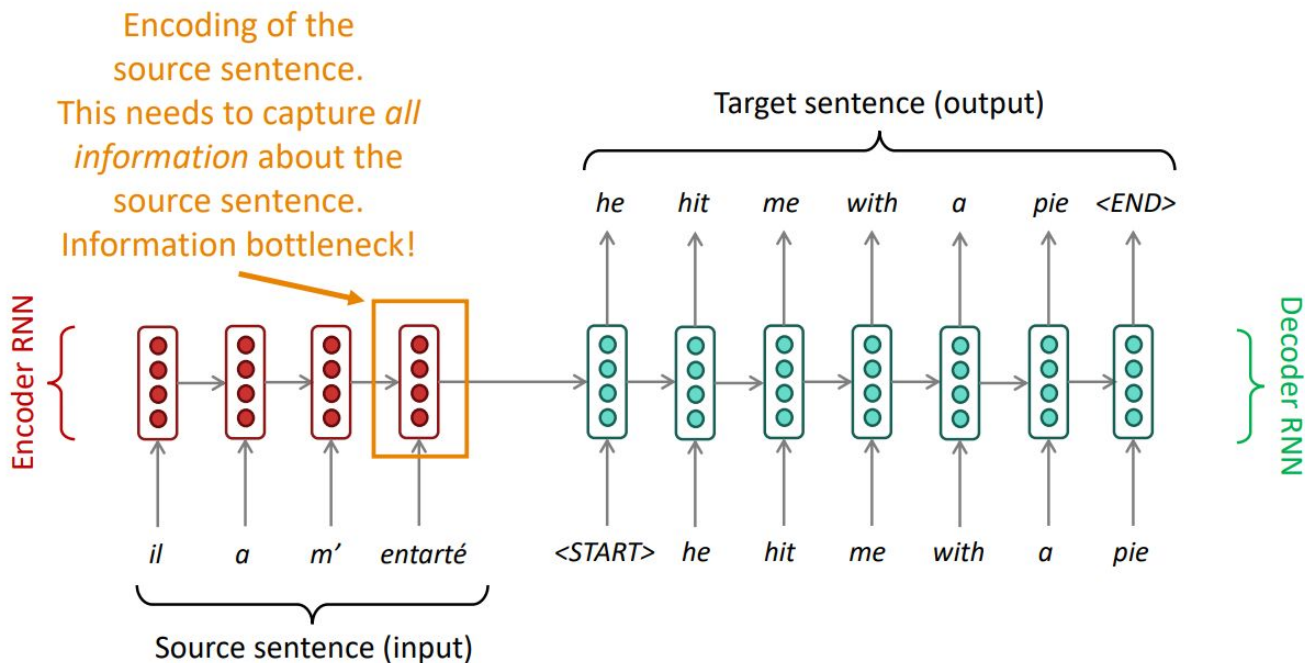


**encoder** processes the input sequence and compresses the information into a context vector of a *fixed length*.

This representation is expected to be a good summary of the meaning of the *whole* source sequence.

!! A critical and apparent disadvantage of this fixed-length context vector design is incapability of remembering long sentences.

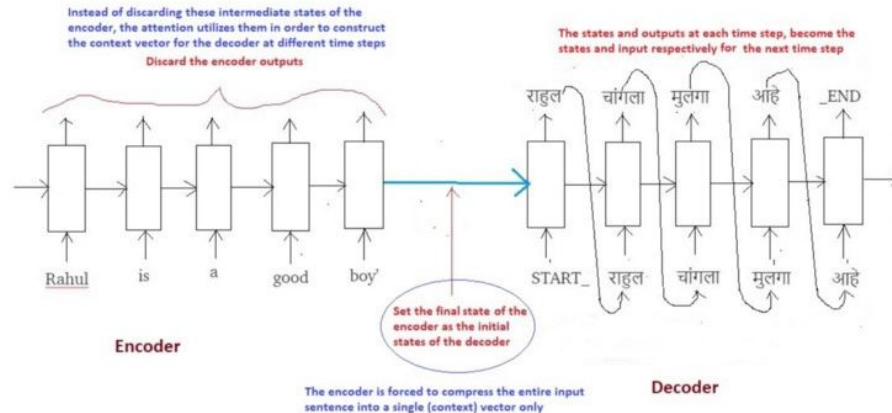
# Sequence-to-sequence: the bottleneck problem



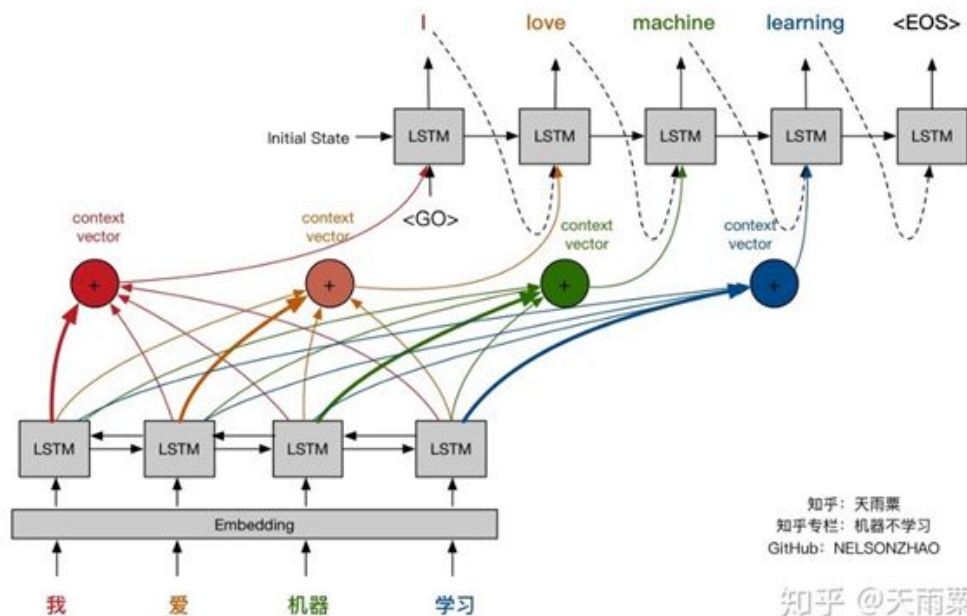


# Born for Translation : Attention


The attention mechanism was born to help memorize long source sentences in neural machine translation. In the traditional **Seq2Seq** model, we discard all the intermediate states of the encoder and **use only its final states (vector) to initialize the decoder**. This technique works good for smaller sequences, however as the length of the sequence increases, a single vector becomes a bottleneck and it gets very difficult to summarize long sequences into a single vector. The central idea behind **Attention** is **not to throw away those intermediate encoder states but to utilize all the states in order to construct the context vectors** required by the decoder to generate the output sequence.




# Encoder-decoder with attention-based mechanism



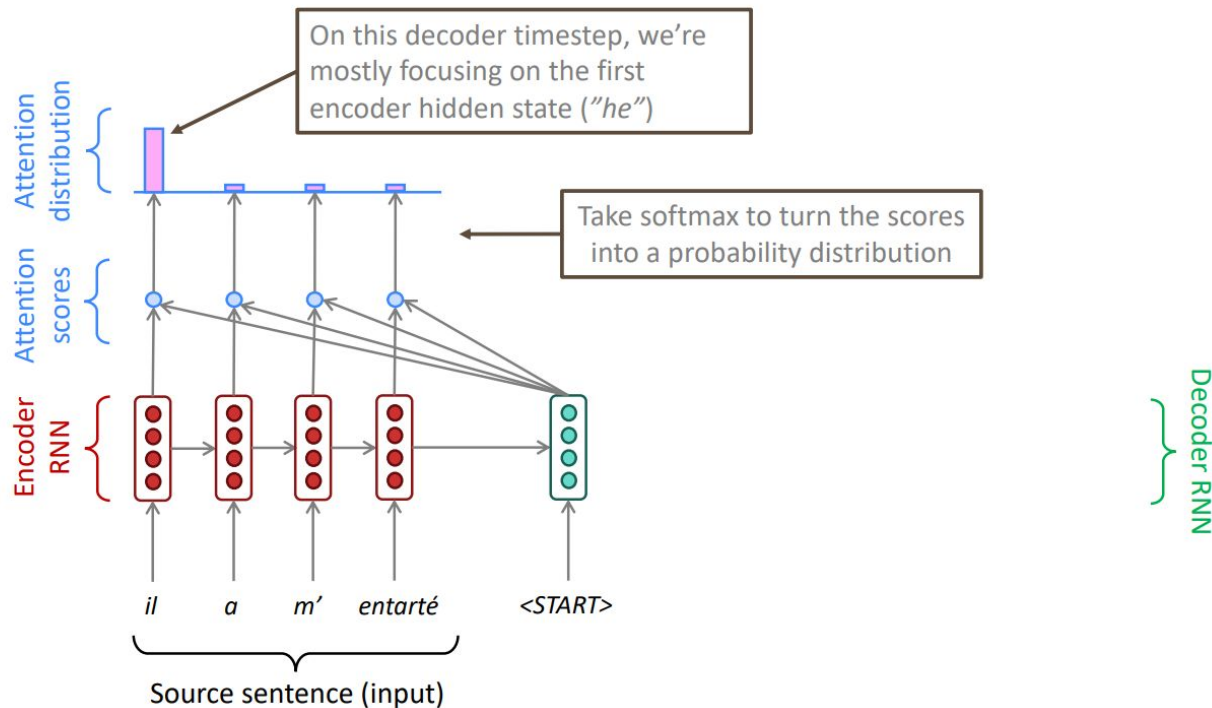
# Sequence to sequence with Attention



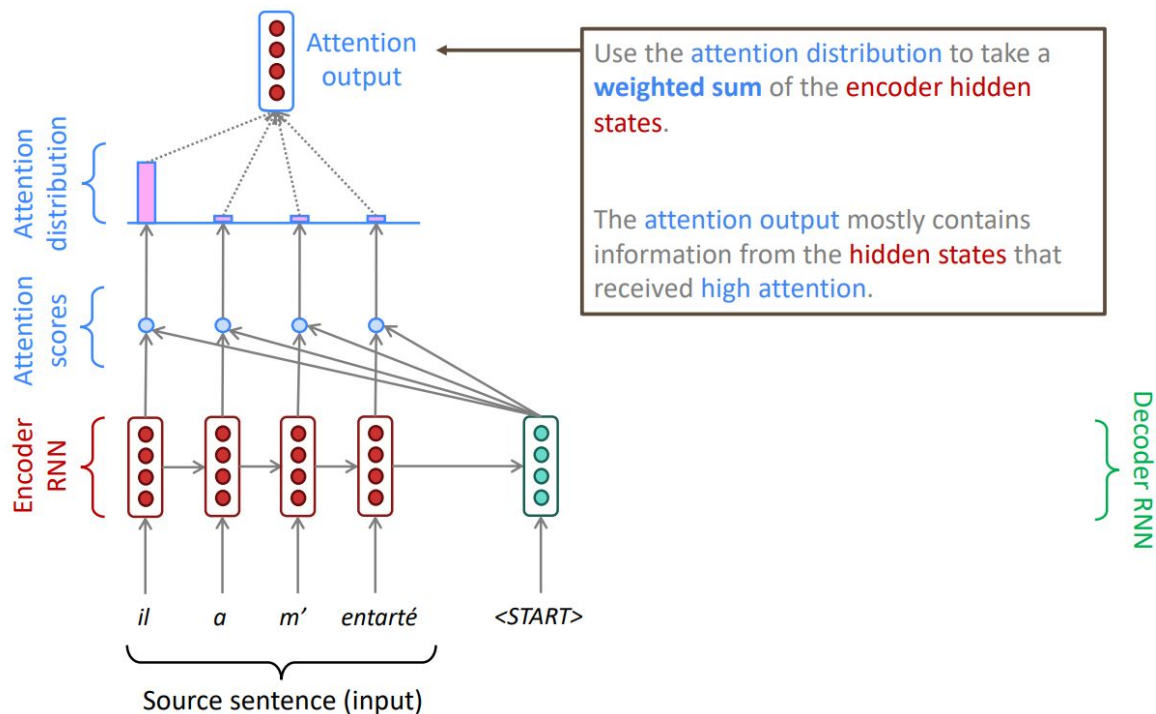
# Sequence to sequence with Attention



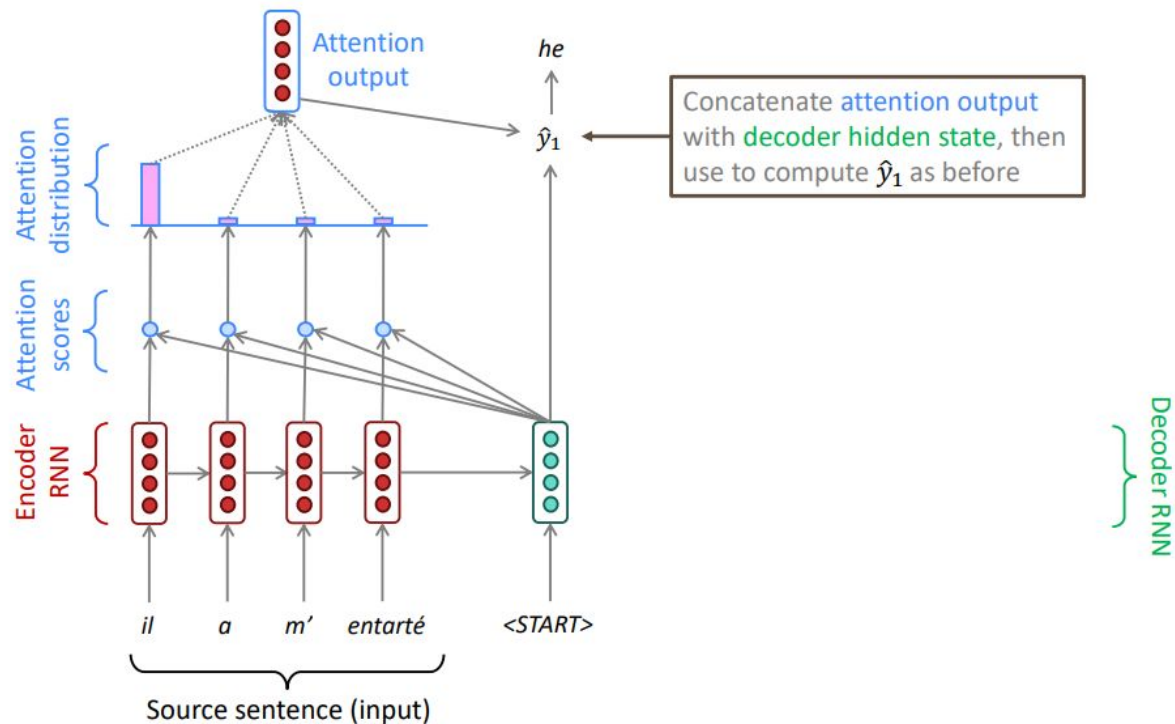
# Sequence to sequence with Attention



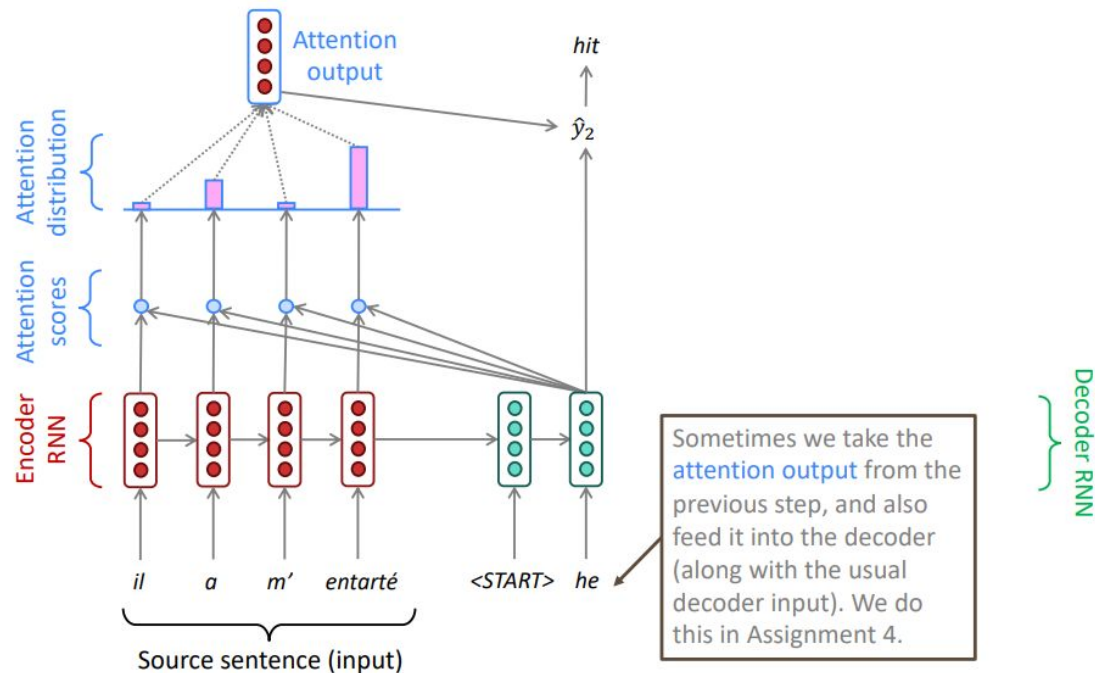
# Sequence to sequence with Attention



# Sequence to sequence with Attention

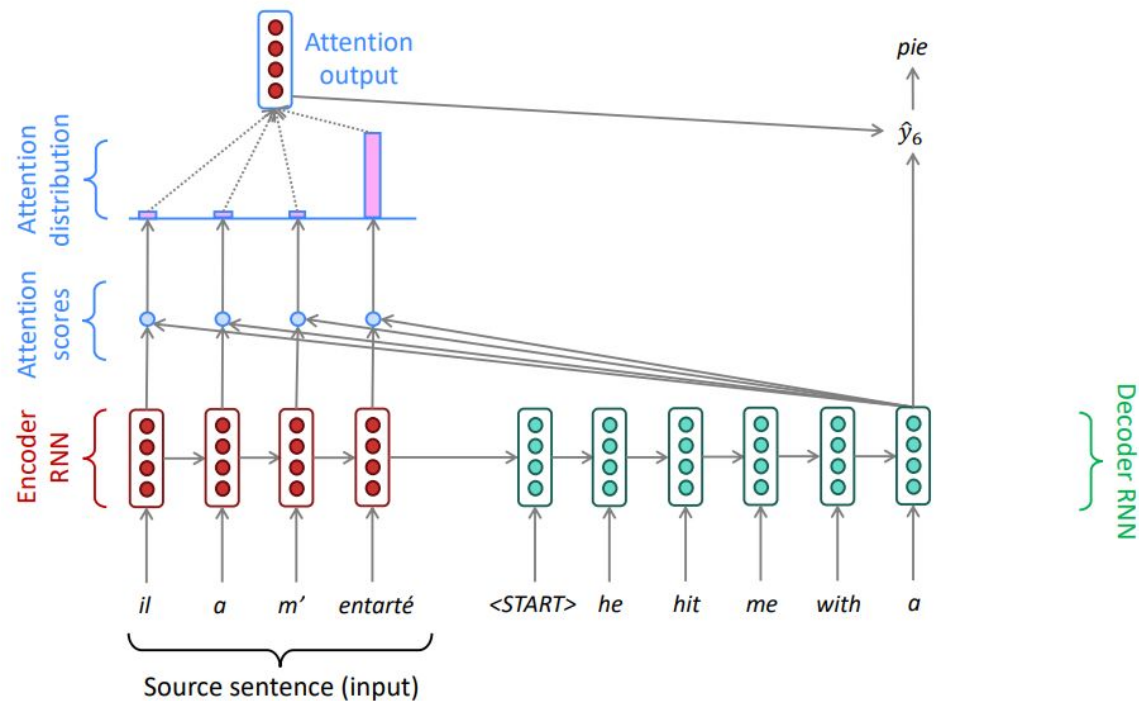


# Sequence to sequence with Attention





# Sequence to sequence with Attention



# Attention: in equations

- We have encoder hidden states  $h_1, \dots, h_N \in \mathbb{R}^h$
- On timestep  $t$ , we have decoder hidden state  $s_t \in \mathbb{R}^h$
- We get the attention scores  $e^t$  for this step:

$$e^t = [s_t^T h_1, \dots, s_t^T h_N] \in \mathbb{R}^N$$

- We take softmax to get the attention distribution  $\alpha^t$  for this step (this is a probability distribution and sums to 1)

$$\alpha^t = \text{softmax}(e^t) \in \mathbb{R}^N$$

- We use  $\alpha^t$  to take a weighted sum of the encoder hidden states to get the attention output  $a_t$

$$a_t = \sum_{i=1}^N \alpha_i^t h_i \in \mathbb{R}^h$$

- Finally we concatenate the attention output  $a_t$  with the decoder hidden state  $s_t$  and proceed as in the non-attention seq2seq model

$$[a_t; s_t] \in \mathbb{R}^{2h}$$

# Variants in Attention

- We have some *values*  $\mathbf{h}_1, \dots, \mathbf{h}_N \in \mathbb{R}^{d_1}$  and a *query*  $\mathbf{s} \in \mathbb{R}^{d_2}$

- Attention always involves:

1. Computing the *attention scores*

$$\mathbf{e} \in \mathbb{R}^N$$

There are  
multiple ways  
to do this

2. Taking softmax to get *attention distribution*  $\alpha$ :

$$\alpha = \text{softmax}(\mathbf{e}) \in \mathbb{R}^N$$

3. Using attention distribution to take weighted sum of values:

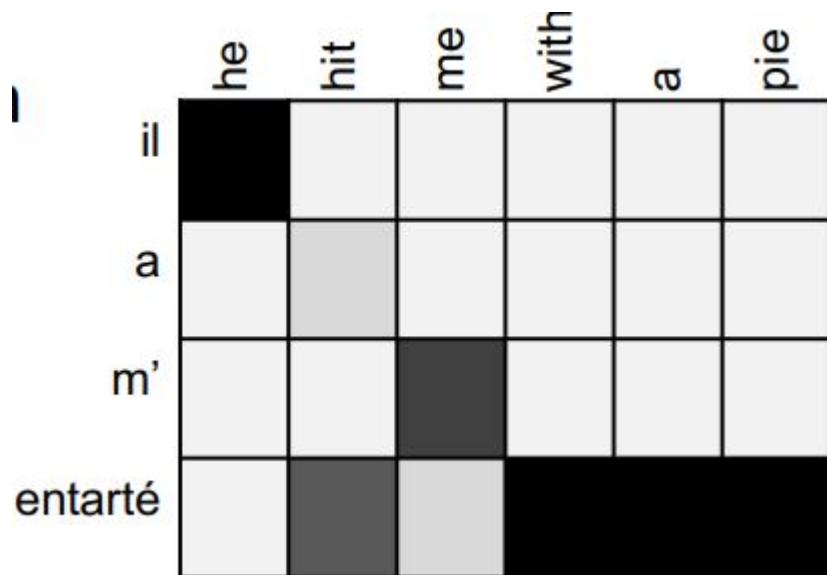
$$\mathbf{a} = \sum_{i=1}^N \alpha_i \mathbf{h}_i \in \mathbb{R}^{d_1}$$

thus obtaining the *attention output*  $\mathbf{a}$  (sometimes called the *context vector*)

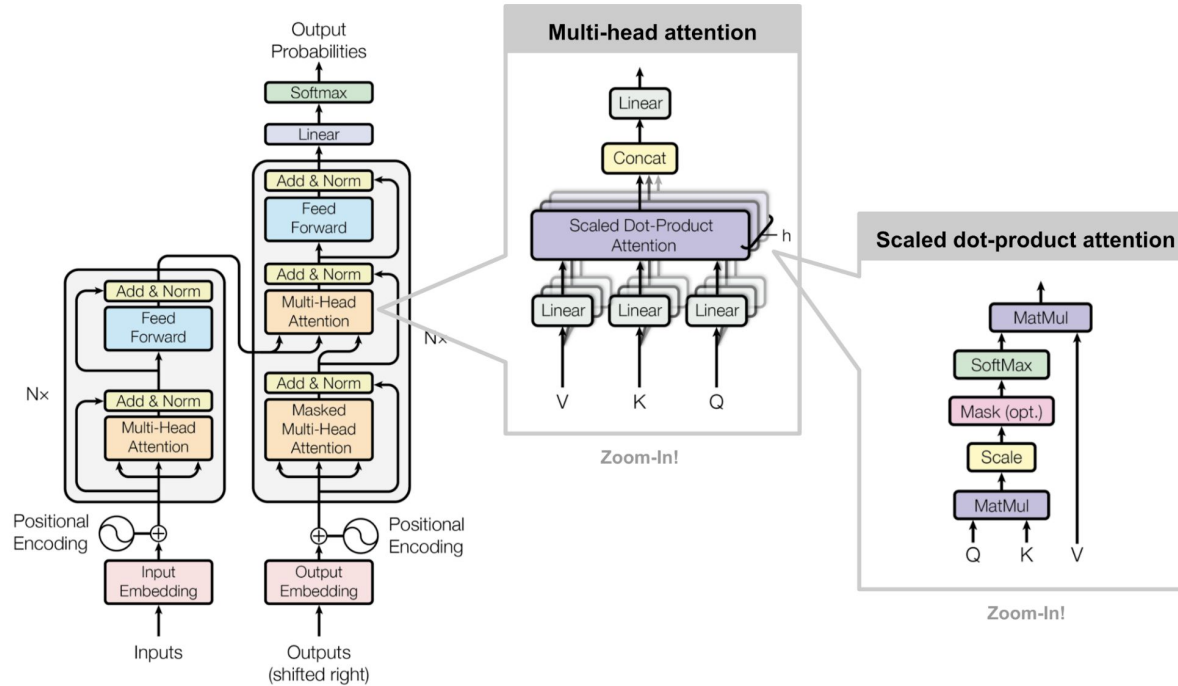
# Types of Score Calculation

Name	Alignment score function	Citation
Content-base attention	$\text{score}(\mathbf{s}_t, \mathbf{h}_i) = \text{cosine}[\mathbf{s}_t, \mathbf{h}_i]$	<u>Graves2014</u>
Additive(*)	$\text{score}(\mathbf{s}_t, \mathbf{h}_i) = \mathbf{v}_a^\top \tanh(\mathbf{W}_a[\mathbf{s}_t; \mathbf{h}_i])$	<u>Bahdanau2015</u>
Location-Base	$\alpha_{t,i} = \text{softmax}(\mathbf{W}_a \mathbf{s}_t)$ Note: This simplifies the softmax alignment to only depend on the target position.	<u>Luong2015</u>
General	$\text{score}(\mathbf{s}_t, \mathbf{h}_i) = \mathbf{s}_t^\top \mathbf{W}_a \mathbf{h}_i$ where $\mathbf{W}_a$ is a trainable weight matrix in the attention layer.	<u>Luong2015</u>
Dot-Product	$\text{score}(\mathbf{s}_t, \mathbf{h}_i) = \mathbf{s}_t^\top \mathbf{h}_i$	<u>Luong2015</u>
Scaled Dot-Product(^)	$\text{score}(\mathbf{s}_t, \mathbf{h}_i) = \frac{\mathbf{s}_t^\top \mathbf{h}_i}{\sqrt{n}}$ Note: very similar to the dot-product attention except for a scaling factor; where n is the dimension of the source hidden state.	<u>Vaswani2017</u>

# Attention Result



# *Attention is All you Need with Transformer Architecture*



# Encoder: Self-Attention

- Step 1: For each word  $x_i$ , calculate its **query**, **key**, and **value**.

$$q_i = W^Q x_i \quad k_i = W^K x_i \quad v_i = W^V x_i$$

- Step 2: Calculate attention score between **query** and **keys**.

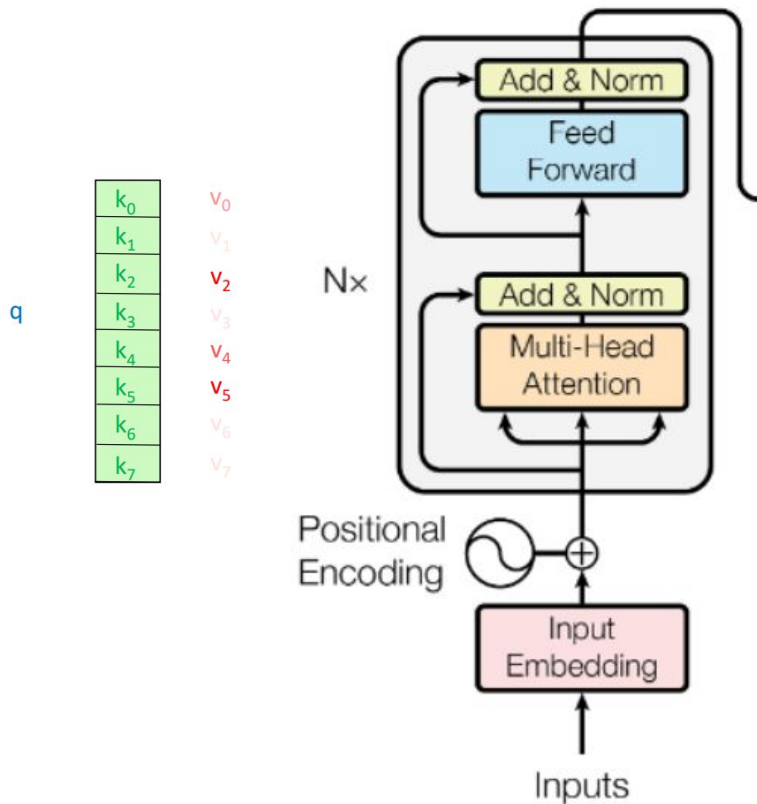
$$e_{ij} = q_i \cdot k_j$$

- Step 3: Take the softmax to normalize attention scores.

$$\alpha_{ij} = \text{softmax}(e_{ij}) = \frac{\exp(e_{ij})}{\sum_k \exp(e_{ik})}$$

- Step 4: Take a weighted sum of **values**.

$$\text{Output}_i = \sum_j \alpha_{ij} v_j$$



# Positional Encoding

Suppose we have an input sequence of length  $L$  and we require the position of the  $k$ th object within this sequence.

$$P(k, 2i) = \sin\left(\frac{k}{n^{2i/d}}\right)$$

$$P(k, 2i + 1) = \cos\left(\frac{k}{n^{2i/d}}\right)$$

$k$ : Position of an object in input sequence,  $0 \leq k < L/2$

$d$ : Dimension of the output embedding space

$P(k, j)$ : Position function for mapping a position  $k$  in the input sequence to index  $(k, j)$  of the positional matrix

$n$ : User defined scalar. Set to 10,000 by the authors of [Attention is all You Need](#).

$i$ : Used for mapping to column indices  $0 \leq i < d/2$ . A single value of  $i$  maps to both sine and cosine functions



# Positional Encoding Example

Sequence	Index of token, $k$	Positional Encoding Matrix with $d=4$ , $n=100$			
		$i=0$	$i=0$	$i=1$	$i=1$
I	0	$P_{00}=\sin(0)$ = 0	$P_{01}=\cos(0)$ = 1	$P_{02}=\sin(0)$ = 0	$P_{03}=\cos(0)$ = 1
am	1	$P_{10}=\sin(1/1)$ = 0.84	$P_{11}=\cos(1/1)$ = 0.54	$P_{12}=\sin(1/10)$ = 0.10	$P_{13}=\cos(1/10)$ = 1.0
a	2	$P_{20}=\sin(2/1)$ = 0.91	$P_{21}=\cos(2/1)$ = -0.42	$P_{22}=\sin(2/10)$ = 0.20	$P_{23}=\cos(2/10)$ = 0.98
Robot	3	$P_{30}=\sin(3/1)$ = 0.14	$P_{31}=\cos(3/1)$ = -0.99	$P_{32}=\sin(3/10)$ = 0.30	$P_{33}=\cos(3/10)$ = 0.96

Positional Encoding Matrix for the sequence 'I am a robot'

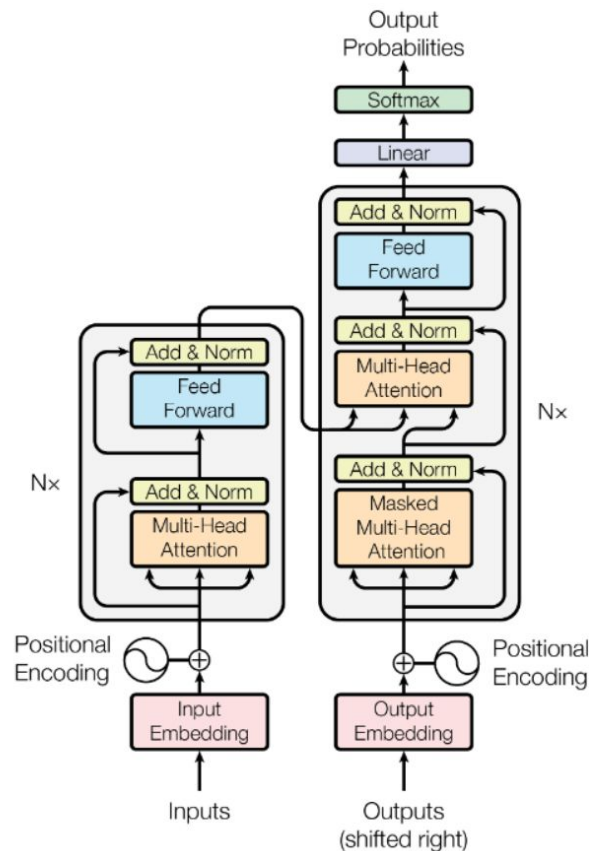

# Masked Multi-Head Attention

At a high-level, we hide (mask) information about future tokens from the model

- To use self-attention in **decoders**, we need to ensure we can't peek at the future.
- At every timestep, we could change the set of **keys** and **queries** to include only past words. (Inefficient!)
- To enable parallelization, we **mask out attention** to future words by setting attention scores to  $-\infty$ .

$$e_{ij} = \begin{cases} q_i^T k_j, & j < i \\ -\infty, & j \geq i \end{cases}$$

For encoding these words



# Image captioning



A person hits a ball with a tennis racket

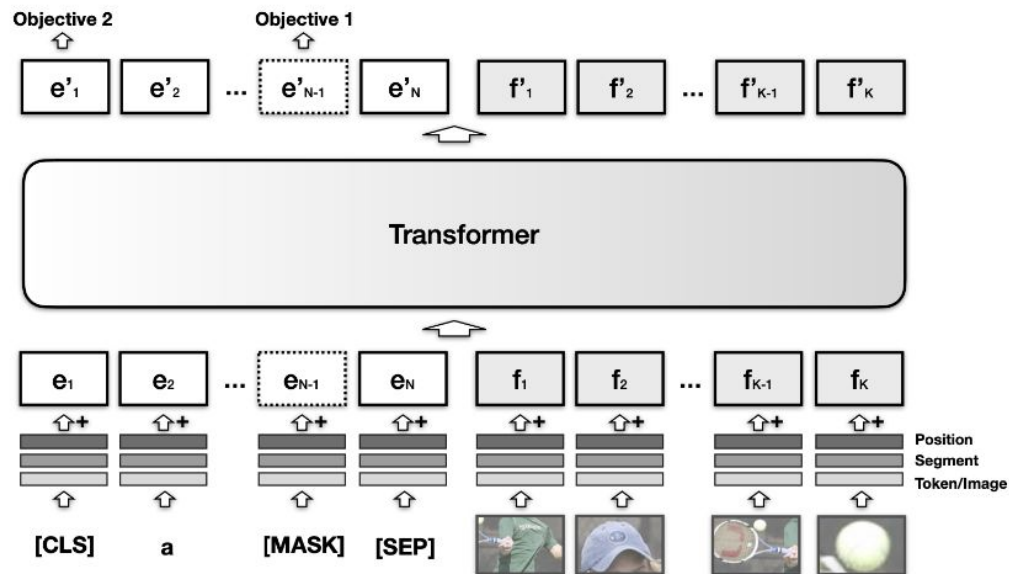


Fig. 1. VisualBERT is trained on the combination of both text and image embeddings. (Image source: [Li et al. 2019](#))

# Image patches to transformer

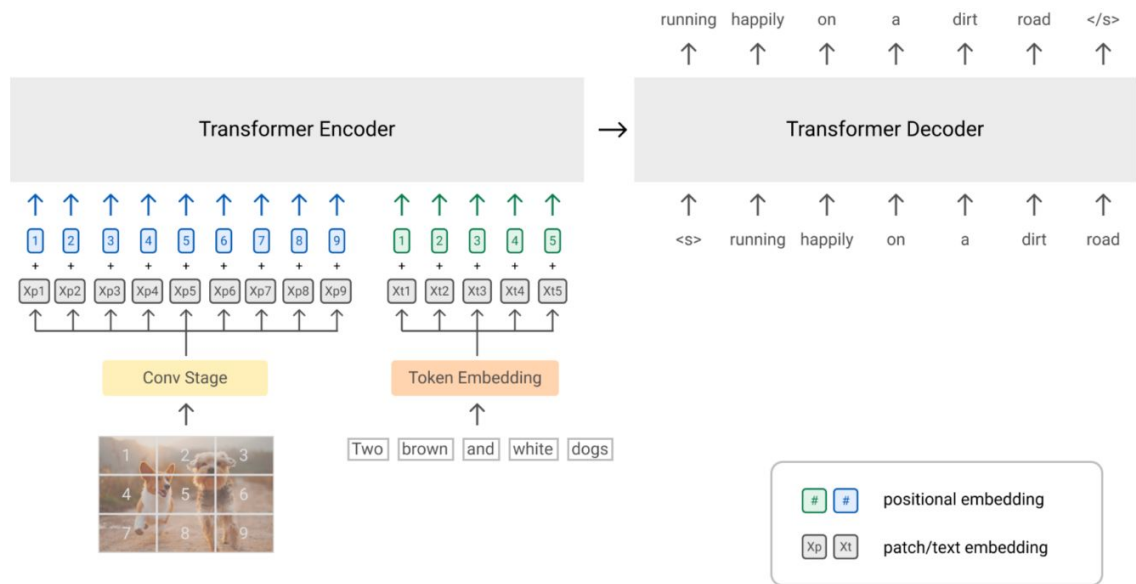


Fig. 3. Training architecture for SimVLM, where the image patches are processed by the cross-attention encoder and the text decoder has causal attention. (Image source: [Wang et al. 2022](#))

# Question answering via Transformer

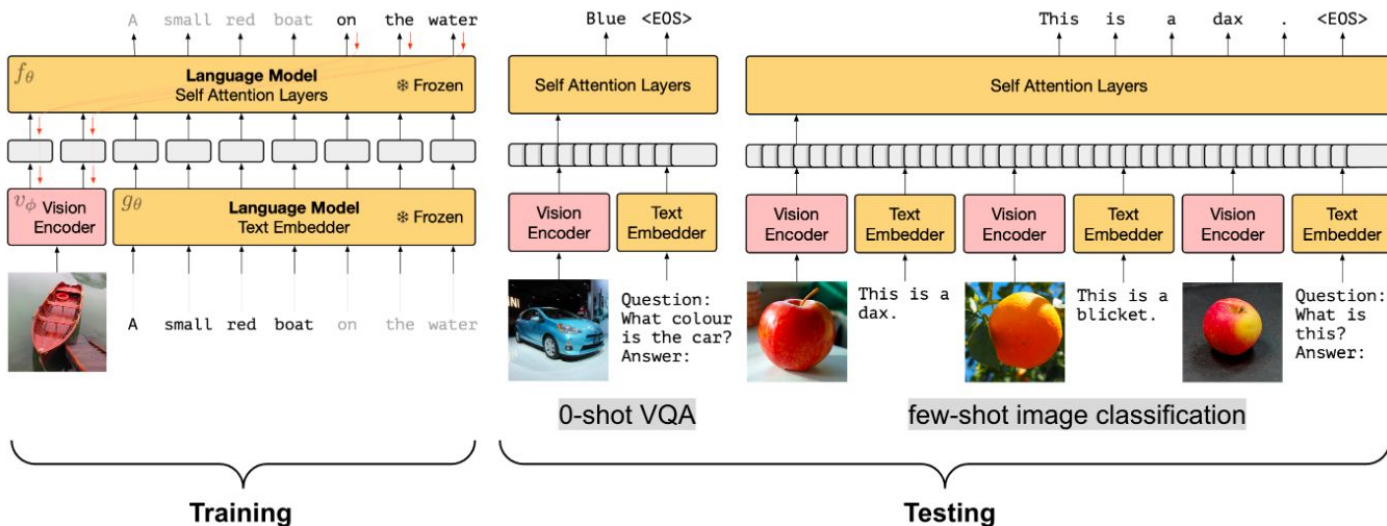


Fig. 6. Illustration of Frozen model (left) training architecture and (right) testing pipeline. (Image source: [Tsimpoukelli et al. 2021](#))

*Thank you for your attention :)*